

Dong-Ki Kim

77 Massachusetts Avenue, Room 31-232C, Cambridge, MA, 02139

☎ (+1) 607-768-6696 | ✉ dkkim93@mit.edu | 🏠 dkkim93.github.io/ | 🎓 Dong-Ki Kim

Education

Massachusetts Institute of Technology

Ph.D. in Aeronautics and Astronautics

Cambridge, MA

Jan. 2020 – Present

- Advisor: Professor Jonathan P. How
- Focus: Multi-Agent Reinforcement Learning
- Cumulative GPA: 5.0 / 5.0

Massachusetts Institute of Technology

S.M. in Aeronautics and Astronautics

Cambridge, MA

Graduated Jan. 2020

- Thesis: “Learning to Teach and Meta-Learning for Sample-Efficient Multiagent Reinforcement Learning”
- Advisor: Professor Jonathan P. How
- Focus: Multi-Agent Reinforcement Learning
- Cumulative GPA: 5.0 / 5.0

Cornell University

B.S. in Electrical and Computer Engineering

Ithaca, NY

Graduated Jan. 2016

- Advisor: Professor Tsuhan Chen
- Focus: Robot Perception
- Highest Honors: *Summa Cum Laude*

Experience

Laboratory for Information and Decision Systems, MIT

Cambridge, MA

Graduate Researcher, Advisor: Professor Jonathan P. How

Sep. 2017 - Present

- Derived new meta-multiagent policy gradient theorem that directly models learning processes of all agents within meta-learning optimization, which enables fast adaptation to new fellow agents across spectrum of mixed incentive, competitive, and cooperative multiagent settings [1, 13].
- Developed safe reinforcement learning framework by learning neural network-based meta-optimizer with projection onto polytope for optimizing objective while satisfying constraints [2, 12].
- Developed peer-to-peer teaching frameworks for enabling agents to learn to teach or share knowledge in cooperative multiagent reinforcement learning settings [3, 4, 5, 14, 15].
- Built attention-based hierarchical reinforcement learning framework that identifies useful latent features across multiple sensory inputs and accelerates in transfer learning tasks [6, 10, 16].
- Led demo preparation of package delivery using multiple drones for annual Boeing visit at MIT. Contributed to collision avoidance algorithm, on-board perception system for classification, and projection system for visualization. [Video]

Air Lab, CMU-Robotics Institute

Pittsburgh, PA

Research Intern, Advisor: Professor Sebastian Scherer

Aug. 2016 - Jul. 2017

- Developed deep multimodal network that improves segmentation robustness to appearance variations (e.g., Summer vs Winter) by combining image and LiDAR sensor data [7, 17].
- Built ROS-based system that estimates terrain roughness from 3D LiDAR sensor data in real-time. [Video]

Robot Intelligence through Perception Lab, TTIC

Chicago, IL

Research Intern, Advisor: Professor Matthew R. Walter

Jan. 2016 - Jul. 2016

- Developed cross-view localization system that estimates vehicle's pose on georeferenced satellite map given sequence of ground-level images [8].
- Improved LSD-SLAM's pose estimation by incorporating ORB-SLAM's pose-graph keyframe constraints.

Advanced Multimedia Processing Lab, Cornell University

Ithaca, NY

Undergraduate Researcher, Advisor: Professor Tsuhan Chen

May. 2014 - Jan. 2016

- Developed indoor localization algorithm based on floor plan and camera [9].
- Built vision-based system that enables drone to navigate indoors autonomously and find specific target [18].

Honor & Award

Outstanding Student Paper Award Honorable Mention for AAAI-19

Jan. 2019

Kwanjeong Education Foundation Scholarship

Sep. 2017 - Present

- Receiving \$30,000 / year for 4-5 years for graduate studies

- Nominated for Cornell University's prestigious award given to top 1% graduating seniors

Skill

Programming Language: Python, C/C++, Matlab

Tools/Library/Software: PyTorch, TensorFlow, Theano, Caffe, Keras, OpenCV, ROS, Point Cloud Library

Publication

Preprint

- [1] **A Policy Gradient Algorithm for Learning to Learn in Multiagent Reinforcement Learning**
Dong-Ki Kim, Miao Liu, Matthew Riemer, Chuangchuang Sun, Marwa Abdulhai, Golnaz Habibi, Sebastian Lopez-Cot, Gerald Tesauro, Jonathan P. How
Under Review as Conference Paper, 2020 [Paper] [Code]
- [2] **FISAR: Forward Invariant Safe Reinforcement Learning with a Deep Neural Network-Based Optimizer**
 Chuangchuang Sun, **Dong-Ki Kim**, Jonathan P. How
Under Review at International Conference on Robotics and Automation (ICRA), 2021 [Paper]

Conference Proceeding

- [3] **Learning Hierarchical Teaching Policies for Cooperative Agents**
Dong-Ki Kim, Miao Liu, Shayegan Omidshafiei, Sebastian Lopez-Cot, Matthew Riemer, Golnaz Habibi, Gerald Tesauro, Sami Mourad, Murray Campbell, Jonathan P. How
International Conference on Autonomous Agents and Multiagent Systems (AAMAS), 2020 [Paper] [WIRED News]
- [4] **Policy Distillation and Value Matching in Multiagent Reinforcement Learning**
 Samir Wadhwan, **Dong-Ki Kim**, Shayegan Omidshafiei, Jonathan P. How
International Conference on Intelligent Robots and Systems (IROS), 2019 [Paper] [Video]
- [5] **Learning to Teach in Cooperative Multiagent Reinforcement Learning**
 Shayegan Omidshafiei, **Dong-Ki Kim**, Miao Liu, Gerald Tesauro, Matthew Riemer, Christopher Amato, Murray Campbell, Jonathan P. How
Association for the Advancement of Artificial Intelligence (AAAI), 2019 [Outstanding Student Paper Honorable Mention] [Paper] [MIT News]
- [6] **Crossmodal Attentive Skill Learner**
 Shayegan Omidshafiei, **Dong-Ki Kim**, Jazon Pazis, Jonathan P. How
International Conference on Autonomous Agents and Multiagent Systems (AAMAS), 2018 [Paper] [Video] [Code]
- [7] **Season-Invariant Semantic Segmentation with A Deep Multimodal Network**
Dong-Ki Kim, Daniel Maturana, Masashi Uenoyama, Sebastian Scherer
Field and Service Robotics (FSR), 2017 [Paper]
- [8] **Satellite Image-based Localization via Learned Embeddings**
Dong-Ki Kim, Matthew R. Walter
International Conference on Robotics and Automation (ICRA), 2017 [Paper] [Video] [NVIDIA News]
- [9] **You Are Here: Mimicking the Human Thinking Process in Reading Floor-Plans**
 Hang Chu, **Dong-Ki Kim**, Tsuhan Chen
International Conference on Computer Vision (ICCV), 2015 [Paper] [Video]

Journal Article

- [10] **Crossmodal Attentive Skill Learner: Learning in Atari and Beyond with Audio-Video Inputs**
Dong-Ki Kim, Shayegan Omidshafiei, Jazon Pazis, Jonathan P. How
Journal of Autonomous Agents and Multiagent Systems (JAAMAS), 2020 [Paper]

Book Chapter

- [11] **Multiagent Reinforcement Learning**

Workshop and Symposium Paper

- [12] **Set-Invariant Constrained Reinforcement Learning with a Meta-Optimizer**
Chuangchuang Sun, **Dong-Ki Kim**, Jonathan P. How
International Conference on Machine Learning (ICML) Workshop, 2020 [Paper]
- [13] **A Policy Gradient Theorem for Learning to Learn in Multiagent Reinforcement Learning**
Dong-Ki Kim, Miao Liu, Matthew Riemer, Golnaz Habibi, Sebastian Lopez-Cot, Samir Wadhwanja, Gerald Tesauro, Jonathan P. How
Association for the Advancement of Artificial Intelligence (AAAI) Spring Symposium, 2020 [Paper]
- [14] **Heterogeneous Knowledge Transfer via Hierarchical Teaching in Cooperative Multiagent Reinforcement Learning**
Dong-Ki Kim, Miao Liu, Shayegan Omidshafiei, Sebastian Lopez-Cot, Matthew Riemer, Gerald Tesauro, Murray Campbell, Golnaz Habibi, Jonathan P. How
Association for the Advancement of Artificial Intelligence (AAAI) Workshop, 2019
- [15] **Learning to Teach in Cooperative Multiagent Reinforcement Learning**
Shayegan Omidshafiei, **Dong-Ki Kim**, Miao Liu, Gerald Tesauro, Matthew Riemer, Christopher Amato, Murray Campbell, Jonathan P. How
International Conference on Machine Learning (ICML) Workshop, 2018
- [16] **Crossmodal Attentive Skill Learner**
Shayegan Omidshafiei, **Dong-Ki Kim**, Jazon Pazis, Jonathan P. How
Neural Information Processing Systems (NeurIPS) Symposium, 2017
- [17] **Online Semantic Mapping for Autonomous Navigation and Scouting**
Daniel Maturana, Sankalp Arora, Po-Wei Chou, **Dong-Ki Kim**, Masashi Uenoyama, Sebastian Scherer
Robotics: Science and Systems (RSS) Workshop, 2017 [Paper]

Technical Report

- [18] **Deep Neural Network for Real-Time Autonomous Indoor Navigation**
Dong-Ki Kim, Tsuhan Chen
arXiv preprint arXiv:1511.04668, 2015 [Paper] [Video]

Academic Service

Program Committee: AAMAS-20

Reviewer of Machine Learning, AI, and Robotics Conferences/Journal: ICLR, AAMAS, ICRA, IROS, IEEE TNNLS, RA-L